

Practical Assignment Report

Practical Assignment 3: Talk Allocation using a CHC Metaheuristic

Authors

Jesús Fernández López & Rafael Carlos Schoettel Vilchez

Subject

Metaheuristics

University of Córdoba

Date

May 19, 2026

Contents

1	Problem Analysis	2
1.1	Context and Motivation	2
1.2	Formal Problem Definition	2
1.2.1	Sets and Parameters	2
1.2.2	Entity Attributes	2
1.2.3	Decision Variable	3
1.2.4	Hard Constraints	3
1.2.5	Soft Constraints (Priorities)	3
2	Design Decisions	4
2.1	Chromosome Representation	4
2.2	Preprocessing: Valid Researchers per Talk	4
3	Fitness Function	5
3.1	Minimisation Approach	5
3.2	Modular Penalty Architecture	5
3.2.1	Hard Constraint Penalties	5
3.2.2	Soft Constraint Penalties	6
3.3	Configuration Dictionary	6
4	Metaheuristic Logic	7
4.1	Overview of CHC	7
4.2	HUX Crossover	7
4.3	Incest Prevention and the Hamming Distance Threshold	8
4.4	Cataclysmic Restart	8
4.5	Repair Operator	9
4.6	Final Solution Guarantee	10
4.7	Elite Set	10
4.8	Algorithm Flow	10
4.9	Gradual Threshold Reset	11
4.10	Stagnation-based Restart	11
4.11	Convergence Behaviour	12
4.12	Parameter Configuration	12
5	Empirical Analysis	13
5.1	Scalability	13
5.2	Fitness Consistency	13
5.3	Feasibility by Instance Size	14
5.4	Impact of Preprocessing	14
5.5	Fitness Breakdown	15

1 Problem Analysis

1.1 Context and Motivation

Every year, to celebrate the *International Day of Women and Girls in Science* (February 11th), the University of Córdoba organises scientific outreach talks at local schools. Schools from the city and the surrounding province send requests: each request states a topic and an educational level for the talk they want. At the same time, researchers volunteer to give talks in their own area of expertise, with some limits on travel and on how many talks they can give.

A coordinator must then match each talk to a suitable researcher. The match must respect a set of hard rules (feasibility) and also follow priority rules that reflect institutional goals. The number of researchers R and the number of talks T change every year, so the algorithm has to handle two different situations: when there are fewer researchers than talks, and when there are more.

1.2 Formal Problem Definition

1.2.1 Sets and Parameters

Symbol	Definition
T	Total number of talks requested by schools
E	Total number of schools
R	Total number of researchers
$\mathcal{T} = \{t_0, \dots, t_{T-1}\}$	Set of requested talks
$\mathcal{S} = \{s_1, \dots, s_E\}$	Set of schools
$\mathcal{R} = \{r_1, \dots, r_R\}$	Set of researchers

1.2.2 Entity Attributes

Each **school** $s \in \mathcal{S}$ is characterised by:

- $\text{loc}(s) \in \{\text{city}, \text{province}\}$ — geographical location.
- $\text{disadv}(s) \in \{0, 1\}$ — whether s is in a disadvantaged area.
- $\text{type}(s) \in \{\text{public}, \text{concerted}, \text{private}\}$ — institution type.
- $\text{new}(s) \in \{0, 1\}$ — whether s participates for the first time.

Each **talk request** $t \in \mathcal{T}$ carries:

- $\text{school}(t) \in \mathcal{S}$ — the requesting school.
- $\text{topic}(t) \in \mathcal{K} \cup \{\text{any}\}$ — required topic ($\mathcal{K}$ is the set of known disciplines; “any” means no restriction).
- $\text{level}(t) \in \mathcal{L}$ — required educational level (preschool, primary, secondary, high school, or vocational training).

Each **researcher** $r \in \mathcal{R}$ is described by:

- $\text{topic}(r) \in \mathcal{K}$ — area of expertise.
- $\text{level}(r) \in \mathcal{L}$ — preferred educational level.
- $\text{travel}(r) \in \{0, 1\}$ — whether r can travel outside the city.
- $\text{newR}(r) \in \{0, 1\}$ — first year participating.
- $\text{prevProv}(r) \in \{0, 1\}$ — gave a talk in the province last year.
- $\text{prevSchool}(r) \in \mathcal{S} \cup \{\emptyset\}$ — school visited last year.
- $\text{maxTalks}(r) \in \{1, 2\}$ — maximum talks r is willing to give.

1.2.3 Decision Variable

This is a **Generalised Assignment Problem** (GAP). The decision variable is a simple mapping:

$$x : \mathcal{T} \rightarrow \mathcal{R} \cup \{-1\} \quad (1)$$

where $x(t) = r$ assigns talk t to researcher r , and $x(t) = -1$ leaves the talk unassigned.

1.2.4 Hard Constraints

HC1. Topic compatibility: $\text{topic}(t) = \text{any}$ or $\text{topic}(r) = \text{topic}(t)$, for all assigned (t, r) .

HC2. Level compatibility: $\text{level}(r) = \text{level}(t)$, for all assigned (t, r) .

HC3. Coverage when $R \geq T$: Every school s must have at least one assigned talk:
 $\exists t: \text{school}(t) = s, x(t) \neq -1$.

HC4. Location feasibility: If $\text{loc}(\text{school}(t)) = \text{province}$ then $\text{travel}(x(t)) = 1$.

HC5. Capacity: $|\{t : x(t) = r\}| \leq \text{maxTalks}(r)$ for all r .

1.2.5 Soft Constraints (Priorities)

When $T > R$ (insufficient researchers), schools are prioritised in the following order:

1. School participates for the first time.
2. School is located in the province.
3. School is in a disadvantaged area.
4. School is a public institution.
5. School is concerted.
6. School is private (lowest priority).

When $R > T$ (surplus researchers), researchers are prioritised:

1. Researcher participates for the first time.
2. Researcher can travel to the province.

Additional historical soft constraints:

- Researcher who was in the province last year should preferably go to the city.
- Researcher should preferably not revisit the same school as last year.

2 Design Decisions

2.1 Chromosome Representation

We encode a solution as a one-dimensional integer array of length T :

$$\mathbf{c} = [c_0, c_1, \dots, c_{T-1}], \quad c_i \in \{-1, 0, \dots, R-1\} \quad (2)$$

Position i of the chromosome corresponds to talk t_i . The value c_i is the *index* of the researcher assigned to that talk (or -1 for unassigned). A direct integer-index encoding was chosen over binary or permutation-based representations because it keeps the chromosome length fixed at T regardless of R , any gene-wise operator (crossover, mutation, repair) applies naturally, and the decoding is immediate: reading position i tells which researcher covers talk i .

2.2 Preprocessing: Valid Researchers per Talk

Before launching the metaheuristic, we pre-compute for each talk t the set $V(t)$ of researchers that can feasibly cover it (stored as a Python dictionary called `valid_researchers_per_talk`):

$$\begin{aligned} V(t) = \{ r \in \mathcal{R} \mid & (\text{topic}(t) = \text{any} \vee \text{topic}(r) = \text{topic}(t)) \\ & \wedge \text{level}(r) = \text{level}(t) \\ & \wedge (\text{loc}(\text{school}(t)) = \text{city} \vee \text{travel}(r) = 1) \} \end{aligned}$$

This **mandatory preprocessing step**:

- Enforces hard constraints HC1–HC4 *before* the search begins, drastically reducing the effective search space.
- Speeds up initialisation: random chromosomes are built by sampling from $V(t)$, producing feasible or near-feasible individuals from generation 0.
- Makes the fitness function lighter: constraint HC1–HC4 violations can be detected in $O(1)$ using the precomputed set.

Talks for which $V(t) = \emptyset$ are inherently infeasible; such genes are set to -1 and are never modified by crossover or mutation.

In the instance generator, a post-processing step ensures every educational level has at least $\lceil T/7.5 \rceil$ researchers. For $T = 30$, this means each of the five levels receives at least 4 researchers, preventing the structural infeasibility that would arise if a level had more talk requests than its researchers could cover.

```

1 def build_valid_researchers_per_talk(talks, researchers, schools):
2     valid = {}
3     for talk in talks:
4         school = schools[talk.school_id]
5         candidates = []
6         for r in researchers.values():
7             # Hard constraint 1: topic match
8             if talk.topic != "any" and r.topic != talk.topic:
9                 continue
10            # Hard constraint 2: level match
11            if r.level != talk.level:
12                continue
13            # Hard constraint 3: travel feasibility
14            if school.location == "province" and not r.can_travel:
15                continue
16            candidates.append(r.researcher_id)
17        valid[talk.talk_id] = candidates
18    return valid

```

Listing 1: Preprocessing: building the valid researchers dictionary

3 Fitness Function

3.1 Minimisation Approach

The fitness function measures the *total penalty* of a chromosome. Following a minimisation paradigm, a perfect solution has $f(\mathbf{c}) = 0$, and each unsatisfied constraint contributes a positive penalty proportional to its weight:

$$f(\mathbf{c}) = \underbrace{\sum_k w_k^{\text{hard}} \cdot \delta_k^{\text{hard}}(\mathbf{c})}_{\text{hard penalties}} + \underbrace{\sum_k w_k^{\text{soft}} \cdot \delta_k^{\text{soft}}(\mathbf{c})}_{\text{soft penalties}} \quad (3)$$

where $\delta_k(\mathbf{c}) \in \{0, 1, \mathbb{Z}^+\}$ counts violations of constraint k .

3.2 Modular Penalty Architecture

Each constraint class has its own function, collecting penalties independently and summing them. This modularity simplifies testing, enables weight tuning, and allows constraints to be enabled or disabled individually.

3.2.1 Hard Constraint Penalties

Table 1: Hard constraint penalties and default weights.

Violation	Description	Default Weight
School without talk ($R \geq T$)	A school receives zero talks when there are enough researchers	1,000,000
Location mismatch	Researcher without travel clearance assigned to a province school	1,000,000
Overallocation	Researcher assigned more talks than their <code>maxTalks</code> limit	1,000,000

Setting hard penalties to 10^6 ensures that any feasible solution — even one that satisfies no soft constraint — outranks any infeasible solution.

3.2.2 Soft Constraint Penalties

Table 2: Soft constraint penalties, their context, and default weights.

Violation	Description	Default Weight
First-year school unserved	School in its first year receives no talk	500
Province school unserved	Province school receives no talk ($T > R$)	400
Disadvantaged area unserved	Disadvantaged school receives no talk	300
Public school unserved	Public institution receives no talk	200
Concerted school unserved	Concerted school receives no talk	100
First-year researcher unused	New researcher not selected ($R > T$)	200
Traveller unused in province	Travelling researcher not used for province ($R > T$)	150
Same school as last year	Researcher revisits their school from last year	300
Same province as last year	Province researcher goes to province again	150
Unassigned talk	No researcher assigned for a requested talk	10

3.3 Configuration Dictionary

All weights are collected in a single Python dict (DEFAULT_CONFIG). This makes the fitness function *fully configurable*: a user can pass a custom dict to `compute_fitness()` to reflect different institutional priorities without modifying any source code.

```

1 DEFAULT_CONFIG = {
2     # Hard constraints
3     "w_school_no_talk": 1_000_000,
4     "w_location_mismatch": 1_000_000,
5     "w_overallocation": 1_000_000,
6     # Soft: unserved schools
7     "w_first_year_school": 500,
8     "w_province_school": 400,
9     "w_disadvantaged": 300,
10    "w_public_institution": 200,
11    "w_concerted": 100,
12    # Soft: unused researchers (R > T regime)
13    "w_first_year_researcher": 200,
14    "w_travel_researcher": 150,
15    # Soft: historical
16    "w_same_school_as_last_year": 300,
17    "w_same_province_as_last_year": 150,
18    # Unassigned talk
19    "w_unassigned_talk": 10,
20 }
```

Listing 2: Configurable weight dictionary

4 Metaheuristic Logic

4.1 Overview of CHC

CHC (Cross-generational elitist selection, Heterogeneous recombination, Cataclysmic mutation) is an evolutionary algorithm proposed by Eshelman (1991). It combines a strong *elitist* strategy (only the best individuals survive) with operators that keep the population diverse:

1. **Cross-generational elitism (C)**: Old and new individuals compete together; only the top N_p survive.
2. **HUX crossover (H)**: A crossover that produces children as different from each other as possible while still mixing both parents.
3. **Cataclysmic restart (C)**: When the population becomes too similar, we restart around the best solution to explore new areas.

Unlike a standard GA, CHC *does not use mutation* during the main search. Instead, it uses incest prevention (two individuals only mate if they are different enough) to avoid premature convergence, and only re-introduces randomness through controlled restarts.

4.2 HUX Crossover

Half-Uniform Crossover (HUX) is the recombination method used in CHC. Given two parents $\mathbf{p}_1$ and $\mathbf{p}_2$:

1. Find all positions where the parents differ: $D = \{i : c_i^{(1)} \neq c_i^{(2)}\}$.
2. Pick exactly half of those positions at random and swap them.
3. The result is two children, each containing a mix of both parents.

$$|D(\mathbf{c}_1, \mathbf{c}_2)| = |D(\mathbf{c}_1, \mathbf{p}_1)| = |D(\mathbf{c}_2, \mathbf{p}_2)| = \left\lfloor \frac{|D(\mathbf{p}_1, \mathbf{p}_2)|}{2} \right\rfloor \quad (4)$$

This equation means each child is equally far from one parent in Hamming distance as from the other parent. This keeps children as diverse as possible while still inheriting genetic material from both parents.

```
1 def hux_crossover(parent1, parent2):
2     diff_positions = [i for i, (a, b) in enumerate(zip(parent1, parent2))
3                     if a != b]
4     if len(diff_positions) < 2:
5         return copy.copy(parent1), copy.copy(parent2)
6     half = len(diff_positions) // 2
7     swap_positions = set(random.sample(diff_positions, half))
8     child1, child2 = copy.copy(parent1), copy.copy(parent2)
9     for pos in swap_positions:
10        child1[pos], child2[pos] = child2[pos], child1[pos]
11    return child1, child2
```

Listing 3: HUX crossover implementation

4.3 Incest Prevention and the Hamming Distance Threshold

CHC uses **incest prevention** to control diversity: two individuals can only mate if their Hamming distance is above a threshold d :

$$d_H(\mathbf{p}_1, \mathbf{p}_2) = \sum_{i=0}^{T-1} \mathbf{1}[c_i^{(1)} \neq c_i^{(2)}] > d \quad (5)$$

The threshold starts at $\lfloor T/4 \rfloor$ and decreases by 1 each time no offspring manages to improve over its parents. This means that, as the population becomes more similar, more pairs are allowed to mate. When the threshold reaches 0, a restart is triggered.

When offspring *do* improve, the threshold rises gradually by $\lfloor T/8 \rfloor$ (rather than resetting all the way to $\lfloor T/4 \rfloor$, as in the original CHC), preserving some of the focused search momentum.

Elitist replacement: A child only enters the population if it is strictly better than at least one of its parents. This guarantees that the best fitness never worsens from one generation to the next.

4.4 Cataclysmic Restart

When d reaches 0 and no offspring can improve over their parents, the population has converged. CHC triggers a **cataclysmic restart** with the following enhancements:

1. **Multiple survivors:** instead of keeping only the single best individual, the restart preserves up to $\lfloor N_p/8 \rfloor$ top unique elite solutions. Each survivor acts as a distinct genetic seed.
2. **Adaptive mutation rate:** starts at $\mu = 50\%$ and increases by 5% for each consecutive restart that fails to improve the best fitness (capped at 80%). This gradually intensifies exploration.
3. Each mutant has $\lfloor \mu \cdot T \rfloor$ genes replaced by a valid researcher drawn from $V(t)$, then passes through the full `repair_chromosome` operator to fix overallocation and fill unassigned positions.
4. The Hamming threshold is reset to $\lfloor T/4 \rfloor$.

Additionally, a **stagnation-based restart** is triggered whenever the best fitness fails to improve for $\max(10, T/2)$ consecutive generations. This helps the algorithm escape plateaus even when the incest threshold has not yet reached zero.

```

1 def cataclysmic_mutation(elites, pop_size, base_rate, stale,
2                           talks, valid_map, r_id_list, researchers):
3     rate = min(base_rate + 0.05 * stale, 0.8)    # adaptive
4     n_mutate = max(1, int(rate * len(talks)))
5     n_surv = min(len(elites), max(1, pop_size // 8))
6     new_pop = [copy.copy(s) for s in elites[:n_surv]]
7     for surv in elites[:n_surv]:
8         for _ in range(...):    # distributed mutants
9             mutant = copy.copy(surv)
10            positions = random.sample(range(T), n_mutate)
11            for pos in positions:
```

```

12         mutant[pos] = repair_gene(pos, valid_map, r_id_list)
13         mutant = repair_chromosome(mutant, r_id_list, researchers,
valid_map)
14         new_pop.append(mutant)
15     return new_pop

```

Listing 4: Cataclysmic restart (simplified)

4.5 Repair Operator

After HUX crossover, the two children might violate the **overallocation** hard constraint: if both parents assign the same researcher to different talks, a child could end up with that researcher assigned to more talks than their `maxTalks` limit allows. The repair operator works in two phases:

1. **Fix overallocation:** detect any researcher with too many talks, randomly pick the extra ones, and try to reassign them to a different valid researcher. If no alternative exists, the talk is left unassigned.
2. **Fill unassigned positions:** scan for genes set to `-1` that have at least one valid researcher with spare capacity, and fill them where possible.

This repair is applied immediately after crossover and also during cataclysmic restart. Reassigning an excess talk may create a new overallocation for the receiving researcher, so the two phases are repeated until the chromosome is fully consistent (up to 10 iterations). This cascade resolution keeps the population close to the feasible region without slowing down the search.

```

1 def repair_chromosome(chromosome, researcher_ids, researchers, valid_map
):
2     chrom = copy.copy(chromosome)
3     for iteration in range(10):
4         changed = False
5         # Phase 1: fix overallocation
6         counts = {}
7         for pos, r_idx in enumerate(chrom):
8             if r_idx != -1:
9                 counts[researcher_ids[r_idx]] = counts.get(..., 0) + 1
10        for r_str, count in counts.items():
11            if count > researchers[r_str].max_talks:
12                # reassign excess talks to other valid researchers
13                ...; changed = True
14        # Phase 2: fill unassigned positions if possible
15        for pos, r_idx in enumerate(chrom):
16            if r_idx == -1:
17                for r_str in valid_map[pos]:
18                    if usage.get(r_str, 0) < researchers[r_str].
max_talks:
19                    chrom[pos] = researcher_ids.index(r_str)
20                    changed = True; break
21        if not changed:
22            break
23    return chrom

```

Listing 5: Iterative repair operator

4.6 Final Solution Guarantee

After the CHC loop finishes, the best chromosome could theoretically carry a remaining overallocation if the valid map is so sparse that no alternative researchers exist. To prevent this, a **forced designation** is applied as a final post-processing step: any researcher assigned to more talks than their `maxTalks` has the excess talks set to -1 (unassigned). This guarantees that the delivered solution never carries a hard penalty of 10^6 , at the cost of a few extra unassigned talks whose individual penalty is only 10. The same cleanup is applied to the elite set, so all displayed alternatives are free of hard violations and directly comparable.

4.7 Elite Set

The algorithm does not return just the single best solution. Instead, it keeps a **diverse elite set** of the top unique individuals found during the entire run. This set is updated each generation and presented to the user at the end, so the coordinator can compare several high-quality alternatives and pick the one that best fits practical needs.

Figure 1 shows the fitness distribution of the final generation from the best-performing run. The top 5 individuals (the elite, in dark blue) are tightly clustered around the optimum, while the remaining population (grey) spans a wider fitness range — evidence that CHC preserves meaningful diversity even after convergence.




Figure 1: Final-population fitness distribution for the best run. The top 5 elite solutions are highlighted in dark blue; the remaining individuals show the diversity maintained by CHC even at convergence.

4.8 Algorithm Flow

Algorithm 6 summarises the complete CHC loop.

```

1 threshold = T // 4
2 stagnation = 0
3 population = initialise_population(pop_size)
4 while generation < max_generations:
5     shuffle(population)
6     pairs = [(pop[i], pop[i+1]) for i in range(0, pop_size-1, 2)]
7     new_individuals = []
8     for p1, p2 in pairs:
9         if hamming(p1, p2) > threshold:
10             c1, c2 = hux_crossover(p1, p2)
11             if fitness(c1) < min(fitness(p1), fitness(p2)):
12                 new_individuals.append(c1)
13             if fitness(c2) < min(fitness(p1), fitness(p2)):
14                 new_individuals.append(c2)
15     if new_individuals:
16         merge and keep top pop_size individuals # cross-
generational elitism
17     else:
18         threshold -= 1 # tighten incest
threshold
19     if threshold <= 0 or stagnation >= max(10, T // 2):
20         population = cataclysmic_restart(elite, rate, stale) #
multi-survivor
21         threshold = T // 4
22         stagnation = 0
23     generation += 1

```

Listing 6: CHC main loop (simplified)

4.9 Gradual Threshold Reset

In the classic CHC algorithm, any improvement resets the incest threshold straight to $\lfloor T/4 \rfloor$. This can be too aggressive: after finding a small improvement, the population is allowed to mate with many very different individuals, losing the local search focus.

We replace the hard reset with a **gradual ramp**:

$$d \leftarrow \min(\lfloor T/4 \rfloor, d + \lfloor T/8 \rfloor) \quad (6)$$

When children improve over their parents, the threshold rises by $\lfloor T/8 \rfloor$ instead of jumping to the ceiling. This keeps the population exploring the current region more thoroughly before reopening the search to distant individuals.

4.10 Stagnation-based Restart

The original CHC only triggers a restart when the incest threshold drops to zero. However, the gradual reset can keep the threshold positive for many generations even though the population has converged and no further improvements are happening.

To handle this, a **stagnation counter** tracks how many consecutive generations pass without a new best fitness. If this counter exceeds $\max(10, \lfloor T/2 \rfloor)$ generations, a cataclysmic restart is triggered regardless of the threshold value. This restart also increments a *stale restart* counter, which gradually increases the mutation rate (explained in Section 4.4) to make successive restarts more disruptive until an improvement is found.

4.11 Convergence Behaviour

Figure 2 shows the average convergence curve over 4 independent runs on a synthetic instance with 18 talks, 5 schools, and 30 researchers ($R > T$). The shaded band represents ± 1 standard deviation. The initial high fitness (around 2900 on average) reflects soft-constraint penalties only — the 2-pass initialisation ensures every school receives at least one talk, so the hard 10^6 penalty is never triggered. Over the first 20–30 generations the algorithm improves the soft priorities, reaching a plateau around 1300. Red dashed lines mark cataclysmic restarts from the best run.




Figure 2: Average convergence over 4 runs on a synthetic instance ($T=18, E=5, R=30$, 150 generations, pop. 40). The shaded band shows $\pm 1\sigma$. Dashed vertical lines mark cataclysmic restarts of the best run.

4.12 Parameter Configuration

Table 3: CHC hyperparameters used in this implementation.

Parameter	Symbol	Default Value
Population size	N_p	50
Maximum generations	G	200
Initial threshold	d_0	$\lfloor T/4 \rfloor$
Restart mutation rate	μ	50% (adaptive up to 80%)

The choice of $\mu = 50\%$ with adaptive scaling provides a stronger exploration baseline than the classic 35% recommendation. When consecutive restarts fail to discover better solutions, the rate increases automatically, making each restart more disruptive until a new improvement is found.

The population size and number of generations scale with the instance: $N_p = \max(200, 5T)$ and $G = \max(500, 20T)$. For the default generated instance ($T = 30$) this gives $N_p = 200$ and $G = 600$, producing run times of 5–30 seconds on consumer hardware.

5 Empirical Analysis

5.1 Scalability

Figure 3 shows execution time as a function of problem size. For a small instance ($T = 15$) the algorithm finishes in about 0.4 seconds. At $T = 60$, with 120 individuals running for 300 generations, the time grows to roughly 5 seconds. The trend is roughly linear in T , since the number of fitness evaluations per generation scales with the population size.




Figure 3: Execution time vs instance size (3 runs per bar, $\pm 1\sigma$).

5.2 Fitness Consistency

When the algorithm is run on realistic instances — without the preprocessing that forces all schools to the city — it faces structurally harder problems (province schools, specific topic requests). Figure 4 shows the final fitness for 5 independent runs on a $T = 60$ realistic instance. Blue bars indicate runs that ended with zero hard penalties (feasible); red bars indicate hard-penalty runs that would still be cleaned up by the forced designation. The tight distribution (with the exception of one outlier) confirms that the algorithm produces consistent solutions regardless of which random seed generates the instance.




Figure 4: Final fitness across 5 runs on a realistic $T=60$ instance. Blue = feasible, red = hard penalty (cleaned post-hoc).

5.3 Feasibility by Instance Size

Figure 5 confirms the algorithm’s robustness: as the instance size grows, the proportion of runs that finish with zero hard penalties increases. At $T = 15$, two out of three runs were feasible; for $T \geq 30$, most runs reached a feasible solution. The forced designation at the end of the algorithm ensures that even the runs with hard penalties are cleaned up before the solution is delivered.




Figure 5: Feasible vs infeasible runs per instance size.

5.4 Impact of Preprocessing

Figure 6 compares two runs of the same instance (seed 42, $T = 30$, pop 80): one realistic (province schools, specific topics) and one after applying the preprocessing that forces every school to the city and accepts any topic. The realistic curve starts with a hard penalty

spike because the travel constraint makes some province schools temporarily unserved; the preprocessed version avoids this entirely. The horizontal dashed line marks the fitness floor after the forced desasignation is applied — the realistic run delivers exactly the same quality as the preprocessed one after this post-processing step.




Figure 6: Convergence comparison: realistic (red) vs preprocessed (blue).

5.5 Fitness Breakdown

When the coordinator inspects a solution, the raw fitness number (1270 in the example shown in Figure 7) is not self-explanatory. The pie chart decomposes it into its constituent penalties: unused-first-year and unused-traveller researchers account for most of the cost (the surplus-researcher regime, $R > T$, is always active); historical repetitions of the same school or province add a smaller fraction; and the tiny unassigned-talk penalty (10 per talk) places the remaining trips. Hard-constraint violations are zero, guaranteed by the forced desasignation.




Figure 7: Penalty composition of the best solution on a $T=30$ instance.

This project tackled the **Talk Allocation Problem**, a type of Generalised Assignment Problem where talks must be matched to researchers subject to several hard rules and a set of priority criteria. The **CHC metaheuristic** was chosen because it combines strong elitism with good diversity control, without relying on hand-tuned mutation probabilities during the main loop. The final implementation was enhanced with adaptive restart mutation, stagnation detection, a two-pass initialisation, and a repair operator — all of which improved population diversity by an order of magnitude (final-population fitness range went from 50 to 550).

The penalty-based fitness function is fully configurable: each constraint has its own weight in a dictionary, so the university can adjust the importance of each priority rule year by year without touching the source code. Hard constraints use a penalty of 10^6 , guaranteeing that any feasible assignment outranks any infeasible one.

A preprocessing step builds $V(t)$ — the set of researchers that can feasibly cover each talk — by filtering on topic, level, and travel. This makes both initialisation and fitness evaluation faster, and lets infeasible talks stay as -1 without wasting search effort.

The algorithm delivers a **guaranteed-feasible solution** by applying an iterative repair operator after each crossover and a forced designation at the end. Even when the problem structure makes overallocation unavoidable, the final chromosome is cleaned to carry only soft penalties. The user also receives a **complete elite set** of the top unique solutions found, displayed as ranked alternatives. On a synthetic instance with $T = 18$, $E = 5$, and $R = 30$, the averaged convergence over several runs starts from a soft-constraint penalty of around 2890 and settles at roughly 1300, showing steady improvement over 150 generations with regular cataclysmic restarts.